

Timeseries Forecasting & Meta-Learning

*William Trouleau, Florian Ravasi,
Sebastien Gay, Gael Grosch,
Dennis Bader, Antoine Madrona*



Unit8™ | SDS **2023**



Florian Ravasi

- MSc Data Science - EPFL
- Data Scientist @ Unit8



Antoine Madrona

- MSc Life Science Engineering - EPFL
- Data Engineer @ Unit8



Sebastian Gay

- BSc Physics and Probability - UniMelb
- Data Engineer @ Unit8



Dennis Bader

- MSc Mechanical Engineering - ETH
- Data Scientist @ Unit8
- Darts Lead



Gael Grosch

- MSc computer science - EPFL
- Data Scientist @ Unit8



William Trouleau

- PhD computer science - EPFL
- Data Scientist @ Unit8



Workshop

Overview

Motivation

- Time series are everywhere
- In Forecasting models need to be handcrafted for each separate applications
- ML is making its entrance in the field
 - Need larger datasets
 - But many problems are still small data problems
- Meta-Learning - definition: applying a model trained on a task (or dataset) onto a different task (or dataset)
 - Open question: Can we use forecasting models learned on certain diverse datasets for other problems?
 - Comparable to what happened to NLP in recent years

Why Forecasting & Meta-learning?

A (brief) history of forecasting techniques...

Soothsayers

ARIMA

Theta

Exponential
Smoothing
+ LSTMs
+ Ensembling

Oracles

Kalman Filters

Exponential
Smoothing

Tea leaves

...

...

Deep learning model
alone, requires no
task-specific tuning

N-BEATS [1]

Promising meta
learning
experiments with
N-BEATS [2]

History

20th century

1998 (M3)

2018 (M4)

2020

2023

Objectives

- Provide an introduction to forecasting with Darts
- Understanding how to forecast multiple time series
- Understanding how ML can be applied to forecasting
- Provide some exposure to meta learning / transfer learning for forecasting
- Build models that get performance close to winning famous M3/M4 forecasting competitions
- Experimenting, playing and having fun!
- ~~Dive in the details of how the models work~~

Pre-requisites

- Basic knowledge of Python for data science
- Basic knowledge of machine learning
- A way to run the provided notebooks, for instance:
 - **[recommended]** Colab environment (<https://research.google.com/colaboratory/>)
 - Your own computer
 - If possible: Nvidia GPU
 - Fresh Python 3.8+ venv

Agenda

1. Workshop overview
2. Introduction to forecasting with Darts
3. Introduction to ML-based forecasting
4. <hands on>
5. Meta-learning overview
6. <hands on>
7. Take home messages



Short intro to Forecasting

Why Forecasting?

Retail



How can we ensure that the consumer demand for every product is met without overspending and creating waste?

Energy



- How much energy should we produce?
- How much energy will we produce given a certain setup?
- What will the electricity price be tomorrow?

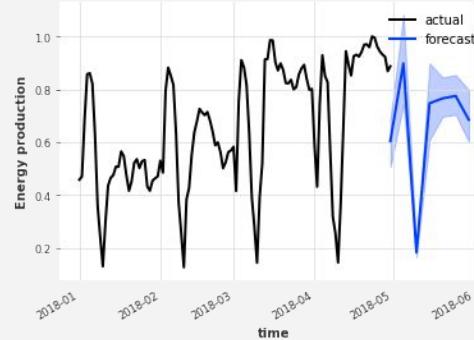
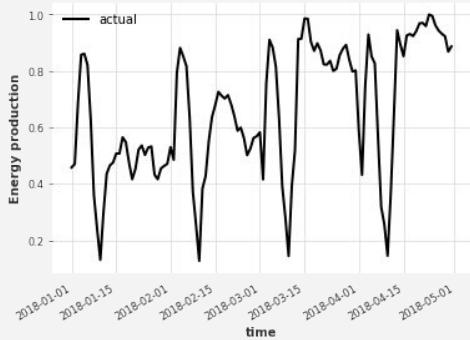
Telecommunications



Will our current infrastructure support the traffic in the near and far future?



Time Series Forecasting



Future = $f(\text{history} + \text{external data})$

Consideration: Do we have a valuable signal in the data?



Darts

Darts is a Python library for **easy manipulation** and **forecasting** of **time series**.



Immutable `TimeSeries` class as basic building block.



Unified `fit()`, `predict()` interface.



Classical models & state-of-the-art ML approaches for forecasting and anomaly detection

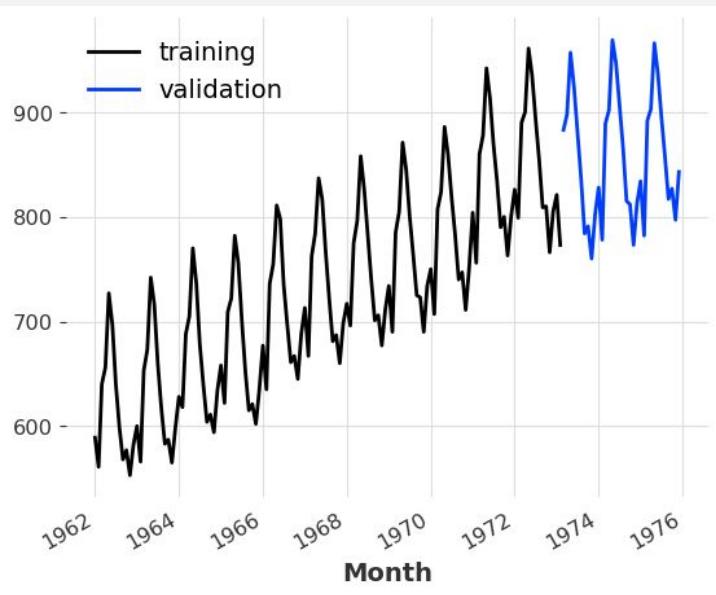


User-friendliness: intuitive API and reasonable defaults.

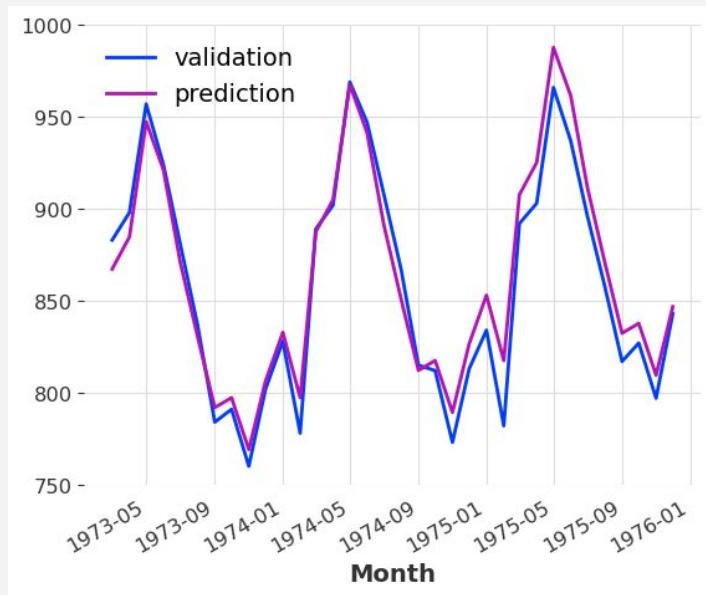


A simple Forecasting example (one series)

Goal



Goal
→



TimeSeries

Processing

Forecasting

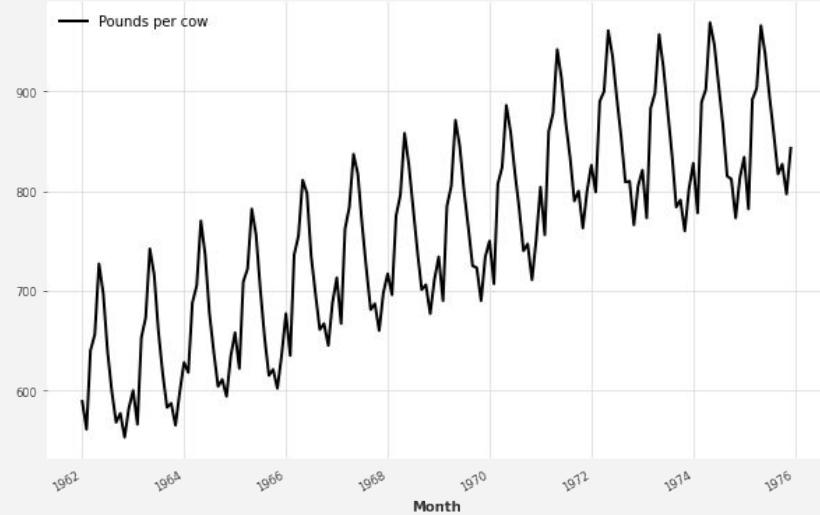
Evaluating

The TimeSeries object



```
from darts import TimeSeries

# Read time series
series = TimeSeries.from_csv('milk.csv', time_col='Month')
series.plot()
```



TimeSeries

Processing

Forecasting

Evaluating

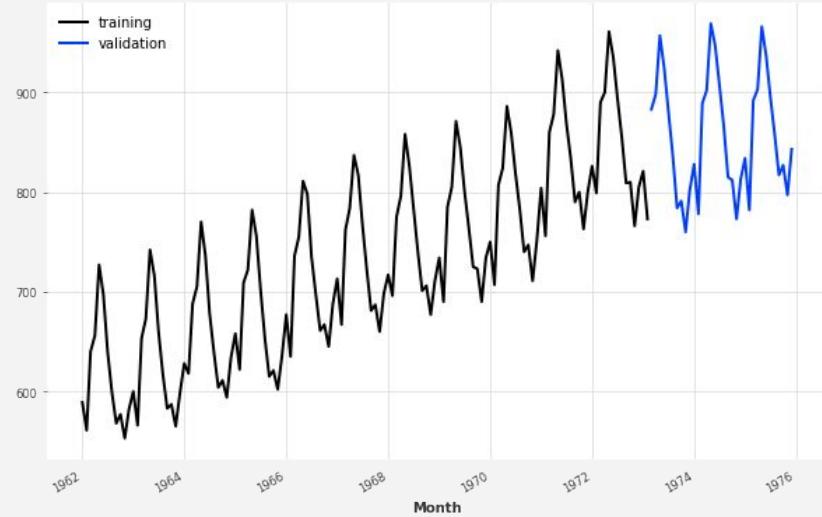
Training / validation split



```
from darts import TimeSeries

# Read time series
series = TimeSeries.from_csv('milk.csv', time_col='Month')
series.plot()

# Split train/val
train, test = series.split_after(0.8)
```



TimeSeries

Processing

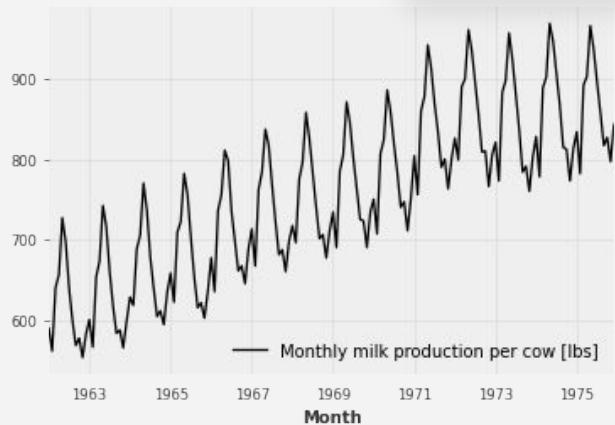
Forecasting

Evaluating

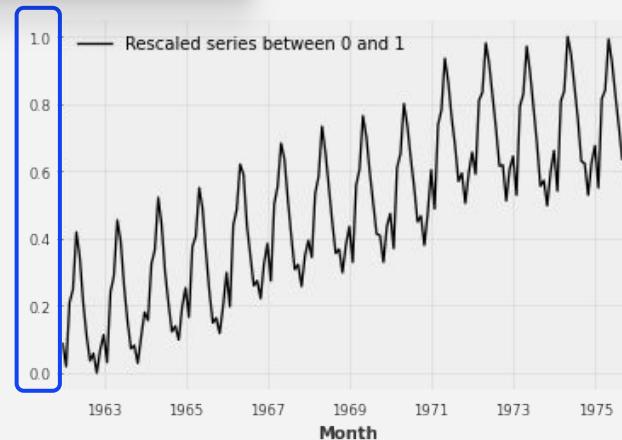
Normalizing TimeSeries



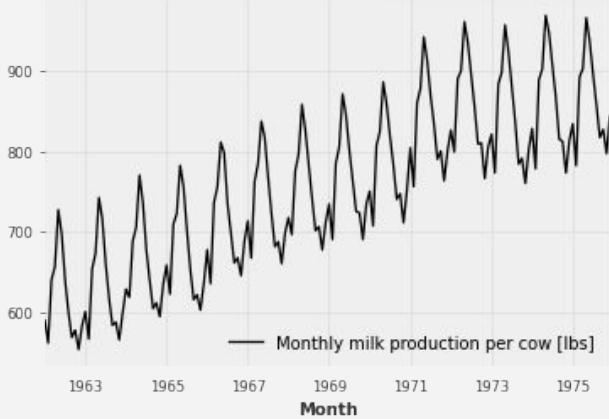
```
scaler = Scaler()  
scaler.fit(series) # Learn min and max from the data  
rescaled = scaler.transform(series)
```



Normalizing

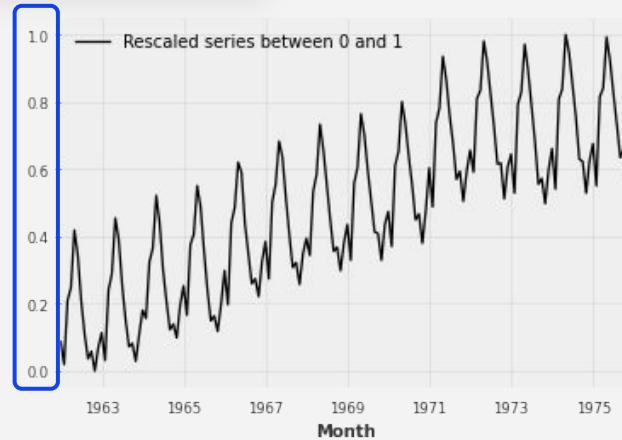


Normalizing TimeSeries



*Re-scaling
back*

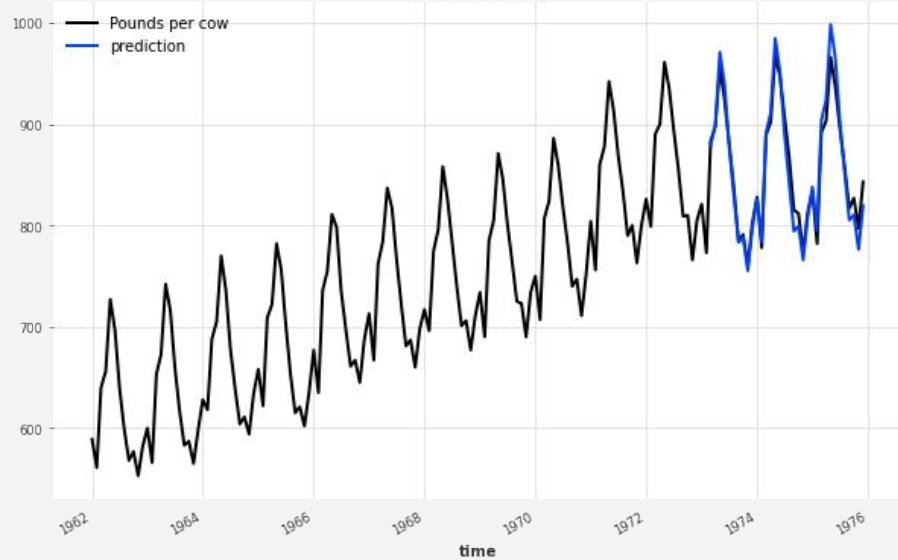
←



Forecasting – Exponential Smoothing



```
from darts.models import ExponentialSmoothing  
  
model = ExponentialSmoothing()  
model.fit(train)  
pred = model.predict(n=len(test))
```



TimeSeries

Processing

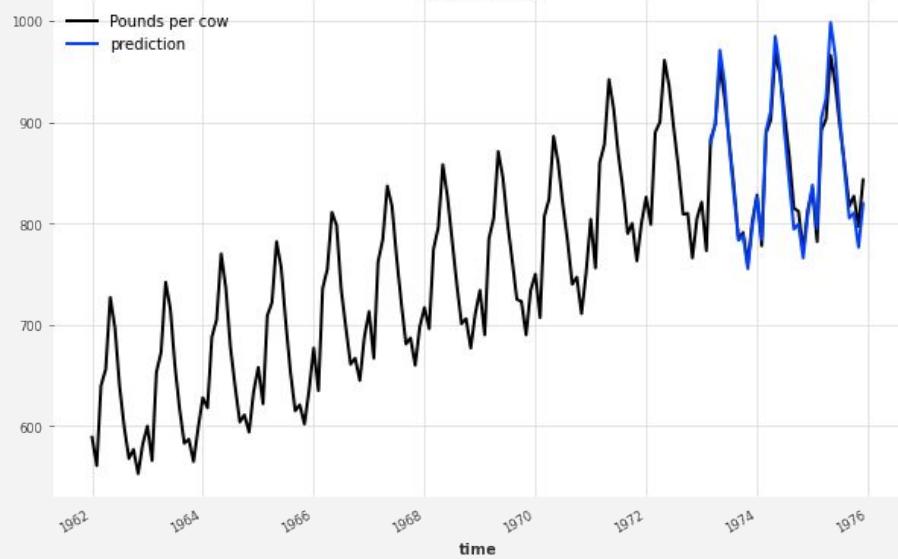
Forecasting

Evaluating

Forecasting – Theta



```
from darts.models import Theta  
  
model = Theta()  
model.fit(train)  
pred = model.predict(n=len(test))
```



TimeSeries

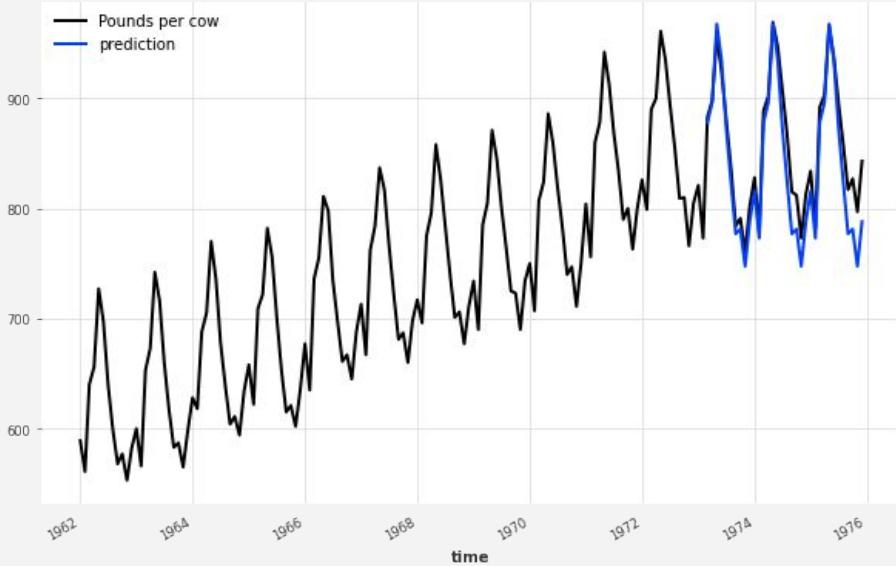
Processing

Forecasting

Evaluating

Specifying parameters

```
from darts.models import Theta  
  
model = Theta(theta=1)  
model.fit(train)  
pred = model.predict(n=len(test))
```



TimeSeries

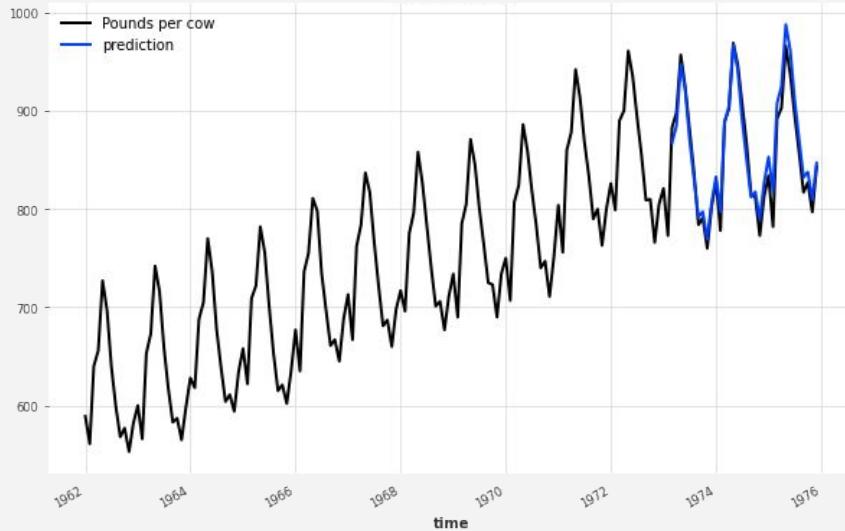
Processing

Forecasting

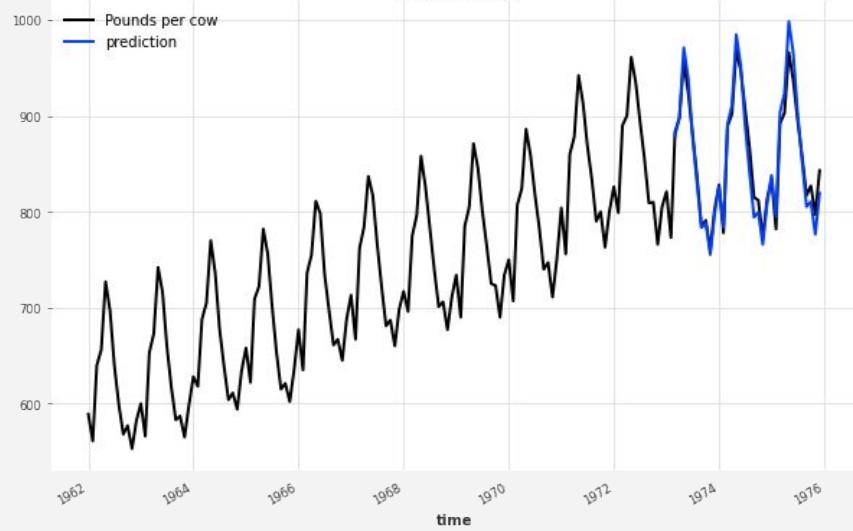
Evaluating

Evaluating predictions – Which one is better?

Exponential Smoothing



Theta



TimeSeries

Processing

Forecasting

Evaluating

Metrics

Many different scores can be computed – Darts lets you import the one you need.



```
from darts.metrics import mape  
  
score = mape(test, pred)
```



```
from darts.metrics import mase  
  
score = mase(test, pred, train)
```

TimeSeries

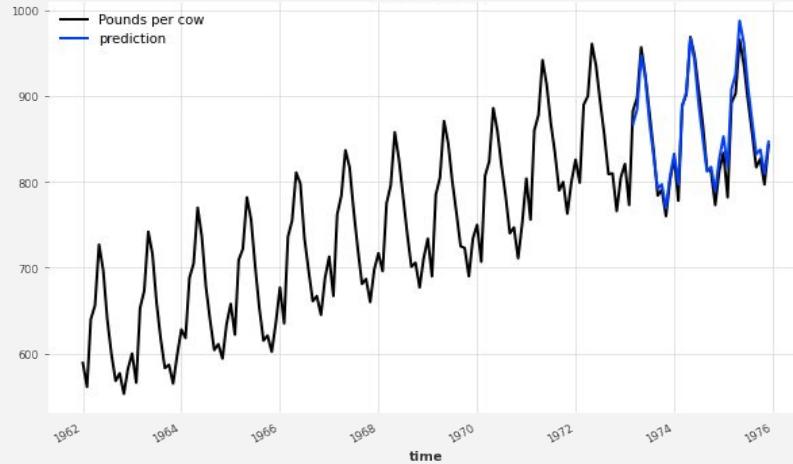
Processing

Forecasting

Evaluating

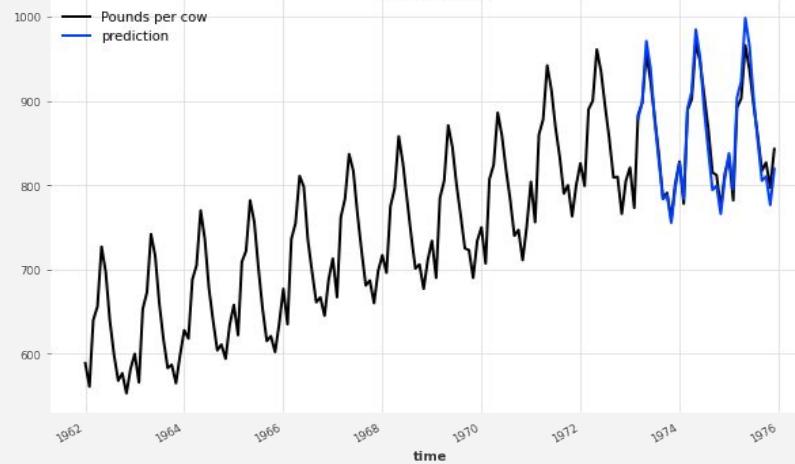
Which one is better?

Exponential Smoothing



MAPE: ~1.39%

Theta



MAPE: ~1.28% 

TimeSeries

Processing

Forecasting

Evaluating

Our Metric: sMAPE

$$\text{MAPE} = 100\% / N \sum_{t=1}^N \frac{|A_t - F_t|}{|A_t|}$$

⬅ Difference between Actual and Forecast
⬅ Normalize by taking Actual

$$\text{sMAPE} = 100\% * 2/N \sum_{t=1}^N \frac{|A_t - F_t|}{|A_t| + |F_t|}$$

⬅ Normalize by taking sum of Actual and Forecast

- sMAPE between 0% and 200%
- MAPE degenerates into positive infinity as soon as any of the actual values is zero



Recap

End-to-end example

```
from darts import TimeSeries
from darts.models import Theta # or others
from darts.metrics import smape
from darts.dataprocessing.transformers import Scaler

# read series, split train/val
series = TimeSeries.from_xxx(...)
train, test = series[:-HORIZON], series[-HORIZON:]

# scale train and val
scaler = Scaler()
train = scaler.fit_transform(train)
test = scaler.transform(test)

# build model with some parameters
model = Theta(...)

# train model
model.fit(train)

# get a prediction
pred: TimeSeries = model.predict(n=HORIZON)

# compute sMAPE error
smape_val = smape(pred, test)
```

Loading, splitting and scaling is taken care of for you in the notebook

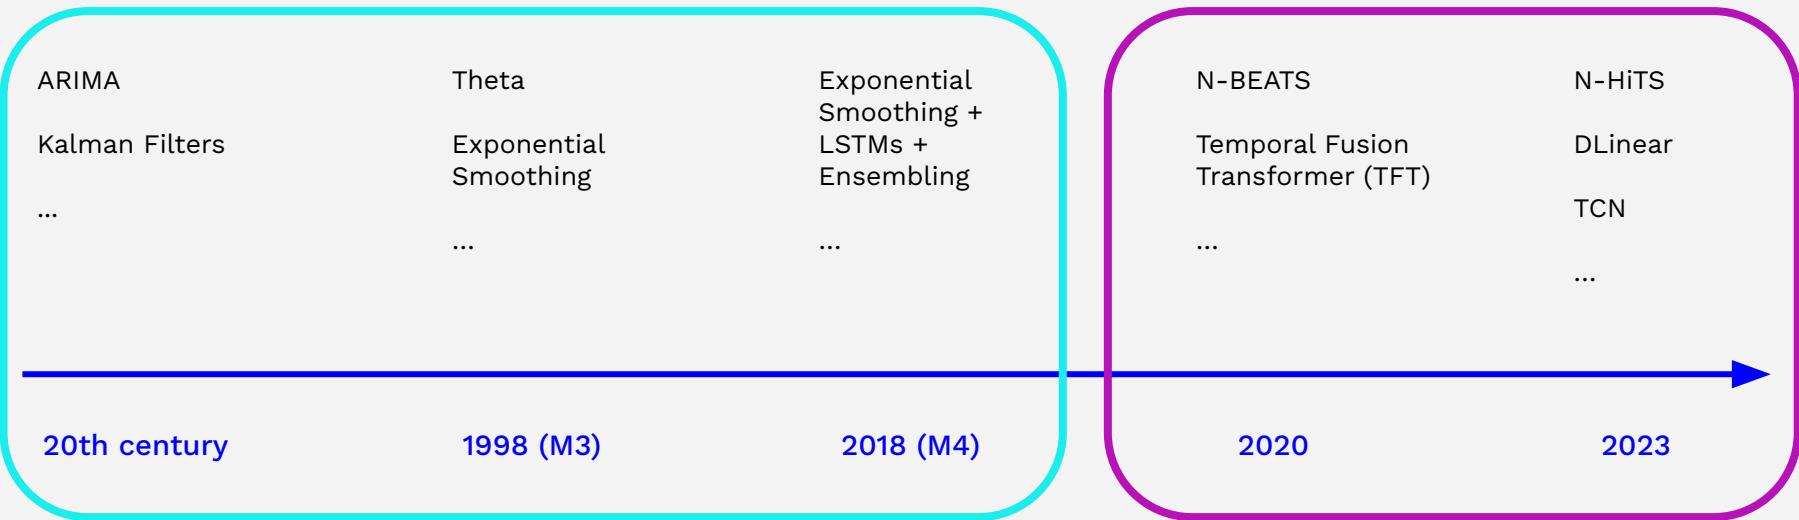
Agenda

1. Workshop overview
2. Introduction to forecasting with Darts
3. **Introduction to ML-based forecasting**
4. <hands on>
5. Meta-learning overview
6. <hands on>
7. Take home messages



Machine Learning for Forecasting

Back to the forecasting techniques history



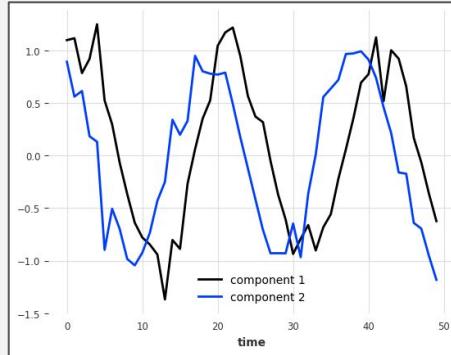
Carefully selected, series specific, parameters
→ **“Local models”, trained with one series**

“Generic” architecture
→ **“Global model”, trained with one or more series**

Multivariate -vs- multi-series

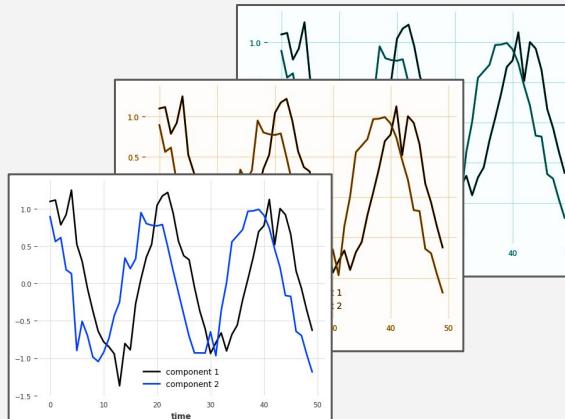
In Darts, one `TimeSeries` can be:

- **Univariate** - i.e. scalar values over time
- **Multivariate** - i.e., vector values over time



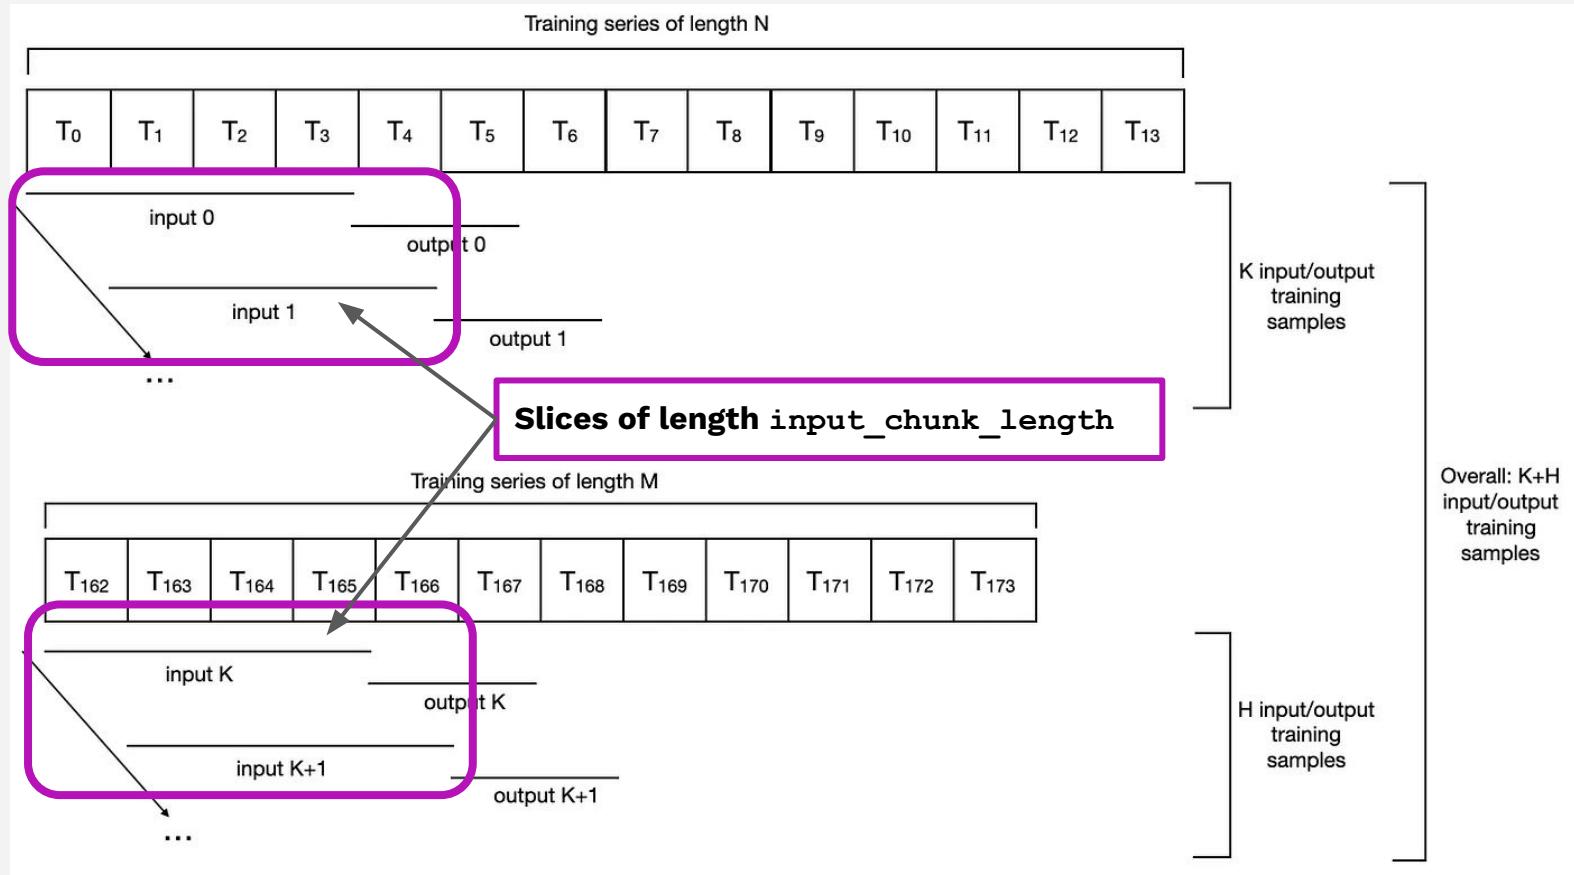
← one multivariate `TimeSeries` with two components

Multiple series : several series, either **univariate** or **multivariate**, with potentially different time axes, scales, trends, ...

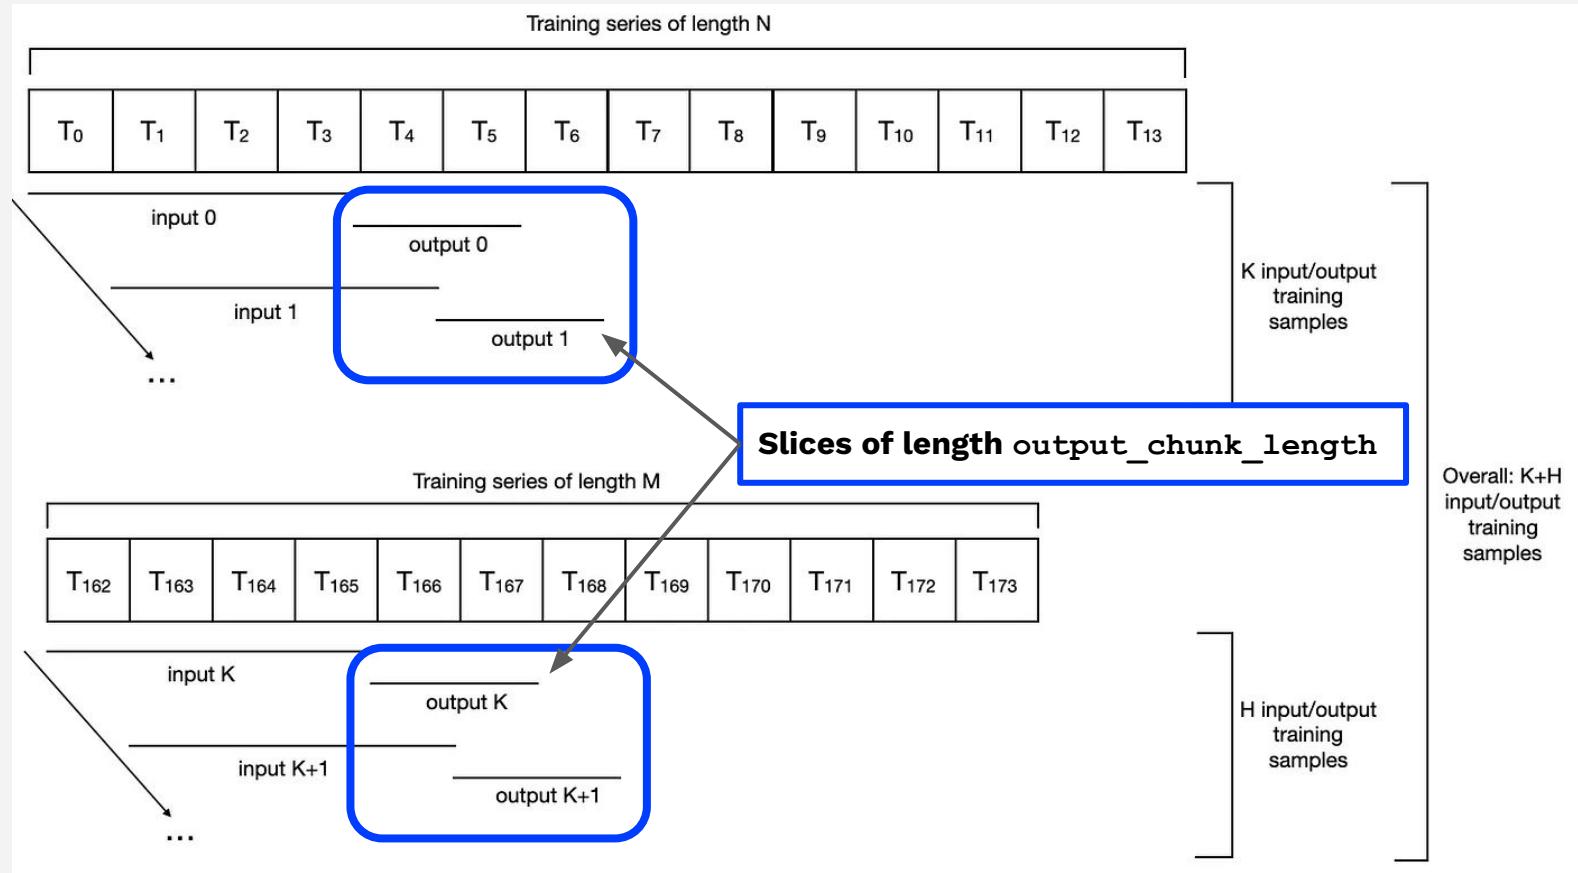


← multiple `TimeSeries`, each corresponding to a different phenomenon

Slicing Multiple Series into a Training Set

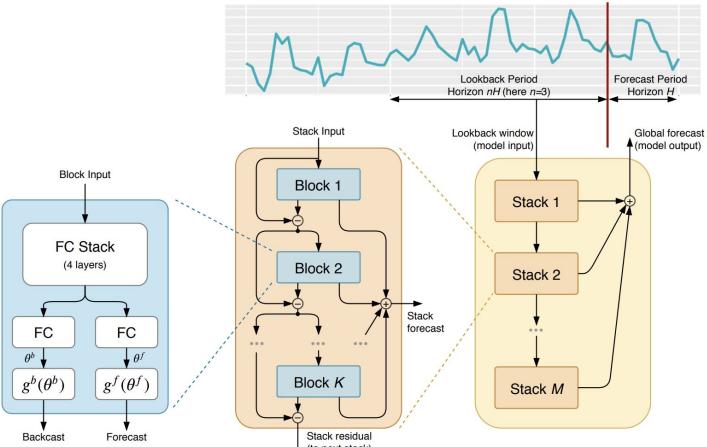


Slicing Multiple Series into a Training Set

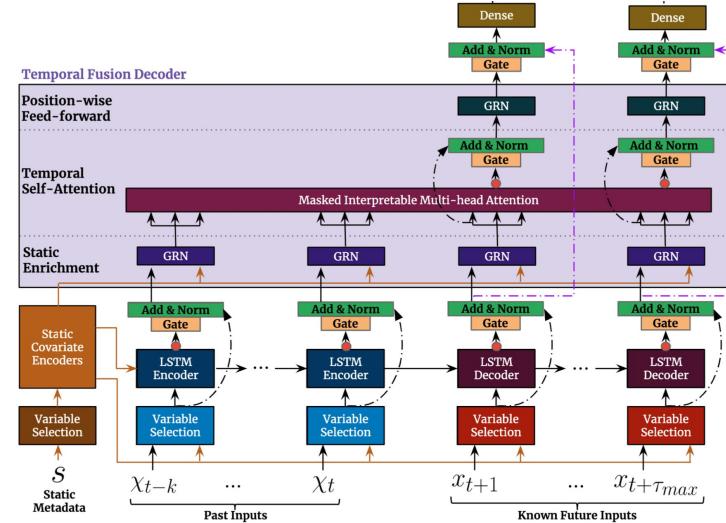


How are these complex architectures trained?

N-BEATS



Temporal Fusion Transformer



Input : Tensor [k+h, input_chunk_length]
Target : Tensor [k+h, output_chunk_length]

Input : Tensor [k+h, input_chunk_length]
Target : Tensor [k+h, output_chunk_length]

In Practice

Fitting and forecasting on collections of series



Agenda

1. Workshop overview
2. Introduction to forecasting with Darts
3. Introduction to ML-based forecasting
4. **<hands on>**
5. Meta-learning overview
6. **<hands on>**
7. Take home messages



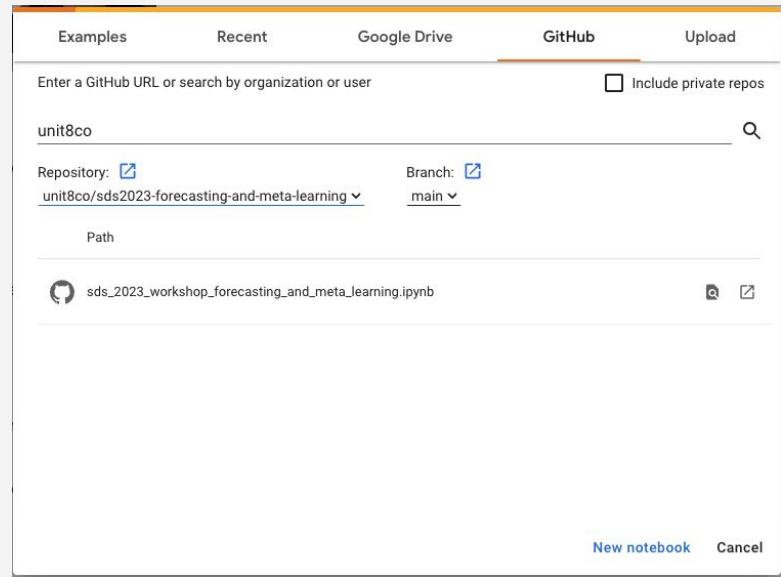
Hands-On Session 1

- Getting Started
- Forecasting multiple series one-by-one
- Global forecasting models to forecast airlines traffic

Setup: Google Colab

We recommend using Google Colab or your local machine for this workshop

- Go to:
<https://colab.research.google.com/>
- "File -> Import notebook -> Github"
- Paste the following link:
<https://github.com/unit8co/sds2023-forecasting-and-meta-learning>
- Load the relevant notebook 



Agenda

1. Workshop overview
2. Introduction to forecasting with Darts
3. Introduction to ML-based forecasting
4. <hands on>
5. **Meta-learning overview**
6. <hands on>
7. **Take home messages**



Meta Learning Overview

Objective: given a set of “Tasks” build a model that can adapt to a new “Task”

“Task”:

- Training dataset
- Validation dataset
- A loss function

“Model”:

- Regular Parameters: Θ
- Meta-Parameters: Φ

Parameters trained using 2 loops:

- **Inner Loop:** given task i create a set of optimised model parameters Θ_i
- **Outer Loop:** based on the Θ_i seen till now update the meta parameters Φ

[1] Oreshkin, Boris N., et al. "Meta-learning framework with applications to zero-shot time-series forecasting". *arXiv preprint arXiv:2002.02887*

Meta Learning in this Workshop

Question: Can we train a model on a “representative” dataset of time series (M4) and then use it to forecasts time series coming from other datasets (e.g. airline traffic)?

Why would we care?

- “Zero-shot” forecasting - no training/engineering/tuning required
- Avoid “cold start” issues on problems with not enough data
- Inference speed



M-Competitions

40 Years of M-Competitions

	Year	# Timeseries	Particularity
M-Competition	1982	1001	- First Competition - Simple vs Complex
M2-Competition	1993	29	- Larger call - 4 actual company data
M3-Competition	2000	3'003	- larger - NN based methods
M4-Competition	2020	100'000	- Much larger set of real life time series - All methods welcomed = ML vs Stat
M5-Competition	2021	42'000	- Kaggle platform - Walmart data
M6-Competition	2022	100	- 50 S&P500 stocks / 50 ETFs - Forecast + investment decision

M4-Competition Dataset

48,000 “diverse” monthly series



Recap

Training on one dataset and testing on another

Agenda

1. Workshop overview
2. Introduction to forecasting with Darts
3. Introduction to ML-based forecasting
4. <hands on>
5. Meta-learning overview
- 6. <hands on>**
7. Take home messages



Hands-On Session 2

- Use a model trained on M4 dataset to forecast other series

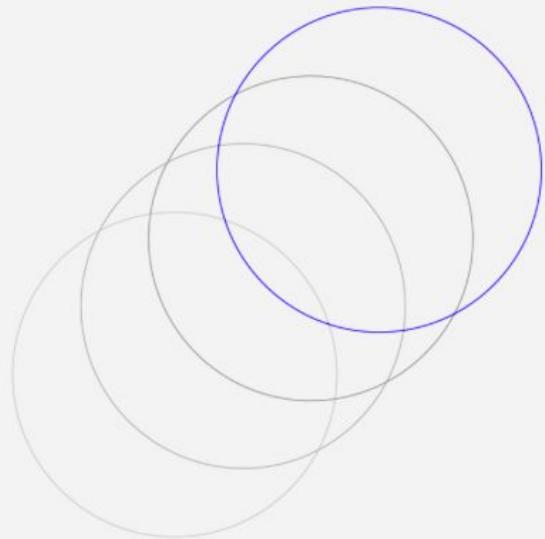
Agenda

1. Workshop overview
2. Introduction to forecasting with Darts
3. Introduction to ML-based forecasting
4. <hands on>
5. Meta-learning overview
6. <hands on>
7. **Take home messages**

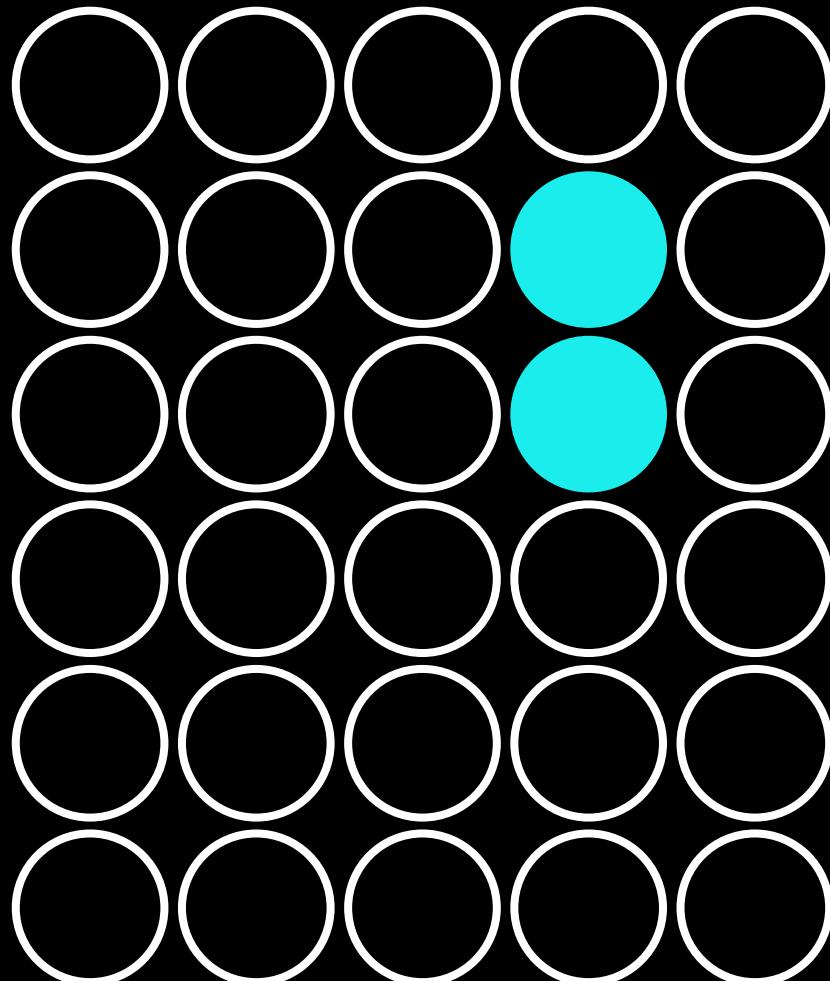


Take home messages

- Darts can be used to very quickly use and compare different models
- “Global models” represent the new way of fitting models on collections of series
- Transferring / meta-learning models from one dataset to another can in some cases reach performance comparable to other models but without requiring training



thank you!



Additional Material

Some definitions

- **Time series** - sequence of values over time
 - **Univariate time series** - a sequence of scalars
 - **Multivariate time series** - a sequence of vectors. Not to be confused with a collection of multiple time series.
- **Forecasting** - making predictions about future values of some series
- **Local forecasting model** - a model trainable on one series to forecast this one series
- **Global forecasting model** - a model trainable on multiple series, to forecast these or other series
- **Meta learning** - applying a model trained on a task (or dataset) onto a different task (or dataset)

Code Snippets

Train like in the NBEATs paper with `HorizonBasedDataset`

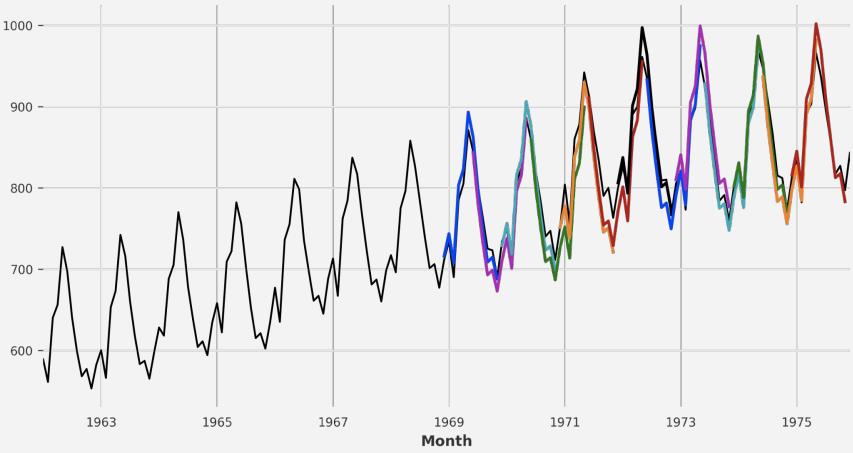
```
● ● ●

train_ds = HorizonBasedDataset(all_train_series, output_chunk_length=12, lh=(1, 3), lookback=7)
val_ds = HorizonBasedDataset(all_val_series, output_chunk_length=12, lh=(1, 3), lookback=7)
print('training set length: {}'.format(len(train_ds)))

_ = nbeats_model.fit_from_dataset(
    train_ds,
    val_dataset=val_ds,
    ...
)
```

Historical simulations using Darts

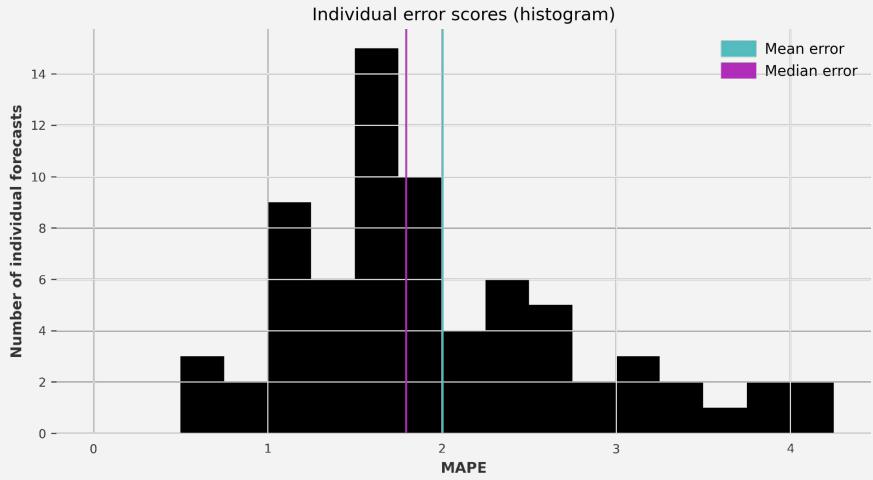
```
● ● ●  
forecasts = model.historical_forecasts(  
    series=series,  
    start=0.5,  
    forecast_horizon=12,  
    stride=6,  
    last_points_only=False  
)
```



Backtesting models on historical data

```
import numpy as np

backtest_errors = model_es.backtest(
    series=series,
    start=0.5,
    forecast_horizon=12,
    stride=6,
    last_points_only=False,
    metric=mape,
    reduction=np.mean
)
```



Early Stopping for Pytorch models

```
● ● ●

early_stopper = EarlyStopping('val_loss', min_delta=0., patience=3, verbose=True)

nbeats_model = NBEATSMModel(
    ...
    pl_trainer_kwarg={"enable_progress_bar": True,
                      "accelerator": "gpu",
                      "gpus": -1,
                      "auto_select_gpus": True,
                      "callbacks": [early_stopper]},
    ...
)

train_ds = HorizonBasedDataset(all_train_series, output_chunk_length=OUT_LEN, lh=LH, lookback=LOOKBACK)
val_ds = HorizonBasedDataset(all_val_series, output_chunk_length=OUT_LEN, lh=LH, lookback=LOOKBACK)

_ = nbeats_model.fit_from_dataset(
    train_ds,
    val_dataset=val_ds,
    num_loader_workers=2,
    epochs=NUM_EPOCHS
)
```

Custom Loss in Pytorch (eg. sMAPE)

```
● ● ●

class SmapeLoss2(nn.Module):
    """
    sMAPE loss as defined in https://robjhyndman.com/hyndsvt/smape/ (Makridakis 1993)
    """
    def __init__(self):
        super().__init__()

    def forward(self, inpt, tgt):
        num = torch.abs(tgt - inpt)
        with torch.no_grad():
            denom = torch.abs(tgt) + torch.abs(inpt)
        return torch.mean(_divide_no_nan(num, denom))

    ...

nbeats_model = NBEATSModel(
    ...
    loss_fn=SmapeLoss2(),
    ...
)
```

Hyperparameter search using Gridsearch

```
parameters = {  
    'theta': [0.5, 1, 1.5, 2, 2.5],  
    'season_mode': [SeasonalityMode.MULTIPLICATIVE,  
                    SeasonalityMode.ADDITIVE]  
}  
  
best_model, best_parameters = Theta.gridsearch(  
    parameters=parameters,  
    series=training,  
    start=0.5,  
    forecast_horizon=12  
)
```



Hyperparameter search using Optuna

```
import optuna

def objective(trial):
    # Configure parameters
    params = {}
    for param in [YOUR_CONTINUOUS_HYPERPARAMETERS]:
        params[param] = trial.suggest_uniform(param, bounds[param][0], bounds[param][1])

    for param in [YOUR_INT_HYPERPARAMETERS]:
        params[param] = trial.suggest_int(param, bounds[param][0], bounds[param][1])

    seasonality = ['additive', 'multiplicative']
    params['seasonality_mode'] = seasonality[trial.suggest_int('seasonality_mode',
                                                               bounds['seasonality_mode'][0],
                                                               bounds['seasonality_mode'][1])]

    .....

    # Train
    m = YOUR_DARTS_MODEL(**params)
    m.fit(train_ts)
    pred = model.predict(val_ts)

    # Return the metric
    return smape(pred, val_ts)

study = optuna.create_study()
study.optimize(objective, n_trials=100)
```

Probabilistic Forecasting in Darts

```
model = TCNModel(likelihood=GaussianLikelihoodModel(),
                  **kwargs)
model.fit(train, covariates)
pred = model.predict(n, covariates=covariates,
                      num_samples=100)
```

You can find all the likelihood available in the darts documentation:

https://unit8co.github.io/darts/generated_api/darts.utils.likelihood_models.html



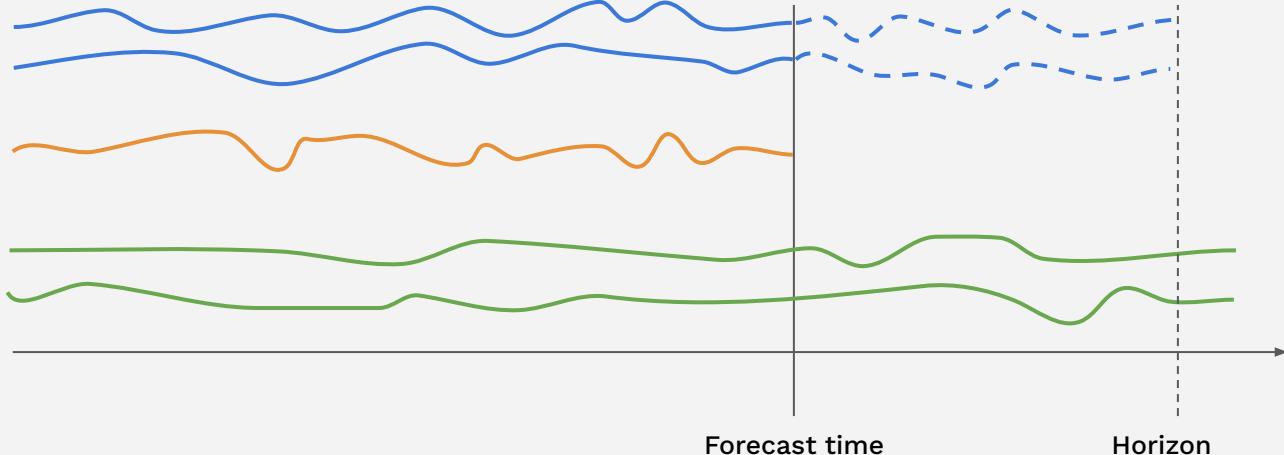
Covariates

Including Past & Future External Data

Target series: what we want to forecast

Past covariates:
Unknown into the future
e.g. measurements

Future covariates:
Known into the future
e.g. calendar, weather forecasts, actions taken



- `fit()` and `predict()` can accept `past_covariates` and/or `future_covariates`, depending on model.
- If `future_covariates` are given, future values will be required at inference time.
- Alignment of covariates with target is automatic.



Training Datasets

[Behind the scenes] Darts Datasets

- “Datasets” classes specify how to slice several series and make a training dataset.
- Out-of-the-box models will build their own datasets (usually simple sequential).
- Datasets wrap around Python Sequence types, so that lazy loading can be used on bigger datasets.

```
● ● ●  
class MyDataset(TrainingDataset):  
    def __init__(  
        self,  
        series: Sequence[TimeSeries],  
        past_covariates: Optional[Sequence[TimeSeries]],  
        ...  
    ):  
  
    def __len__(self,):  
        return len(self.series)  
  
    def __getitem__(self, i: int):  
        # Return i-th (input, output) example
```

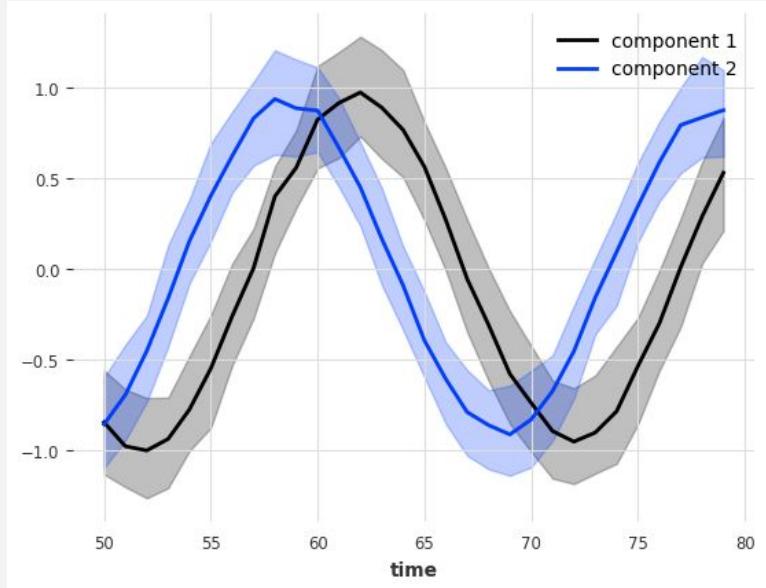
Probabilistic Forecasting

Probabilistic Forecasting



```
from darts.models import KalmanForecaster  
  
model = KalmanForecaster(dim_x=2)  
model.fit(series)  
pred = model.predict(n=30, num_samples=100)  
  
pred.plot()
```

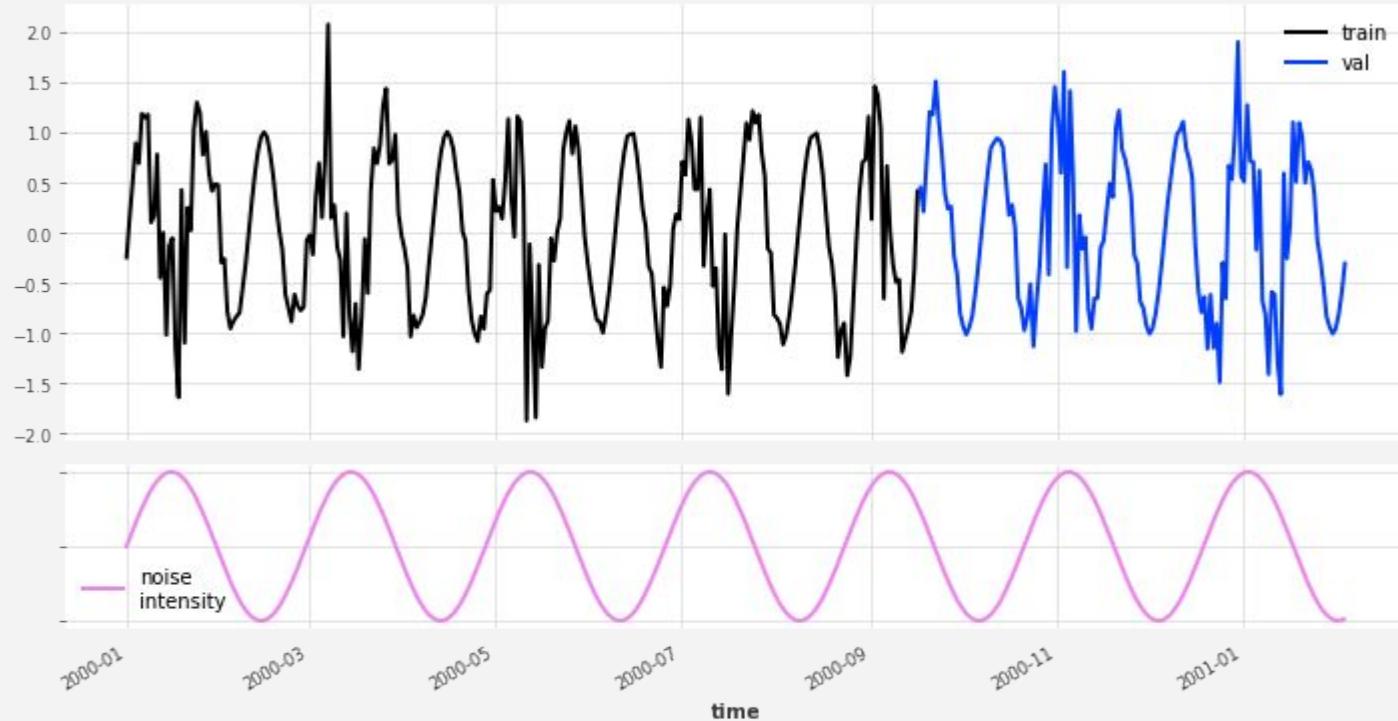
one probabilistic **TimeSeries**:



TimeSeries can be **deterministic** or **stochastic**.
Some models provide probabilistic forecasts.



Probabilistic Forecasts - Example with a noisy sine wave

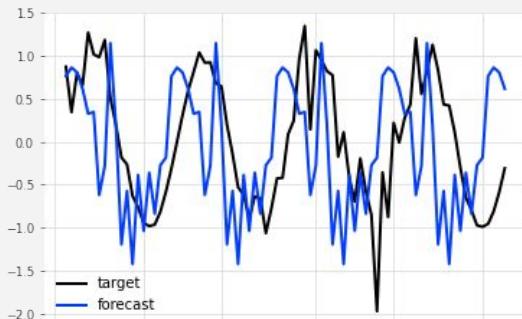


Capturing Series Stochasticity

Attempt 1

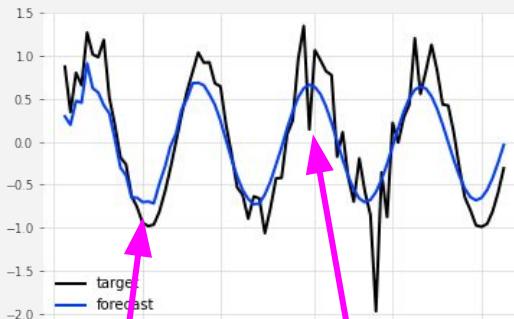


```
model = NaiveSeasonal(seasonal_period)
model.fit(train)
pred = model.predict(n)
```



Attempt 2


```
model = ARIMA(seasonal_period, 0, 0)
model.fit(train)
pred = model.predict(n)
```



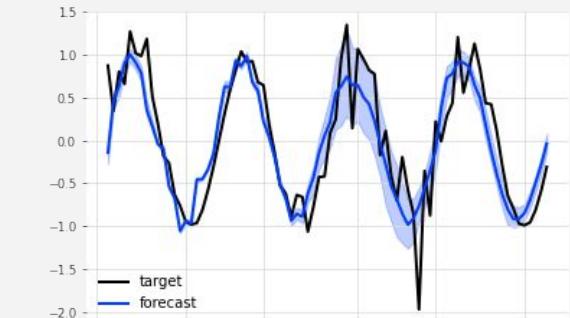
“safe” forecast

uncertain forecast

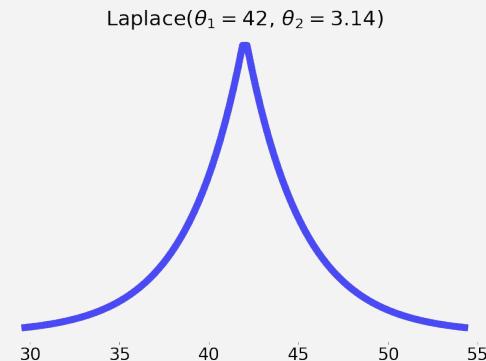
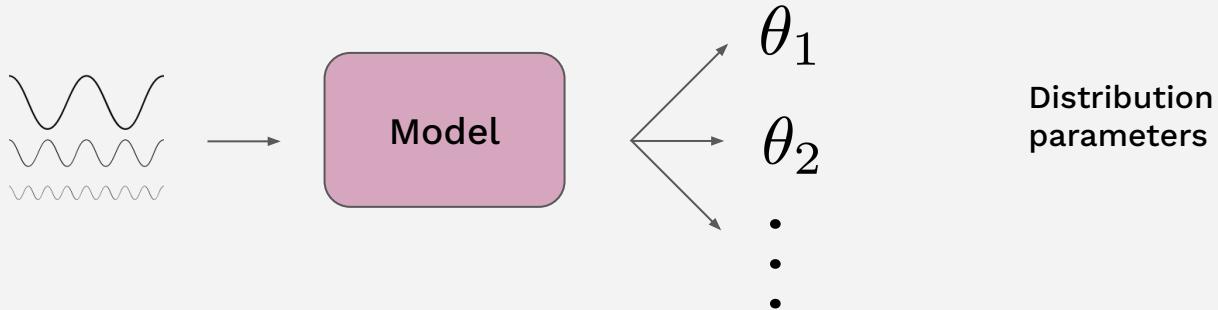
Attempt 3



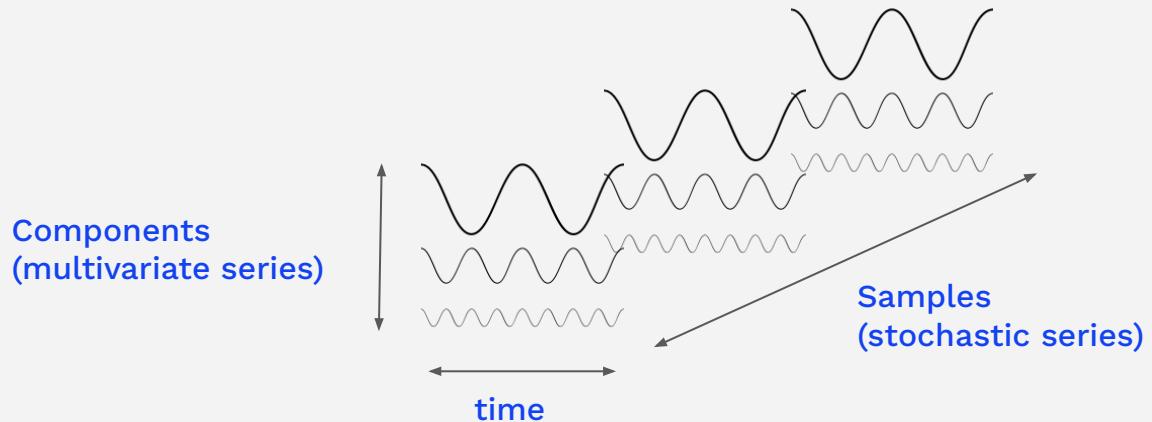
```
model = TCNModel(likelihood=GaussianLikelihood(),
                  **kwargs)
model.fit(train, past_covariates=noise_intensity)
pred = model.predict(n,
                      past_covariates=noise_intensity,
                      num_samples=100)
```



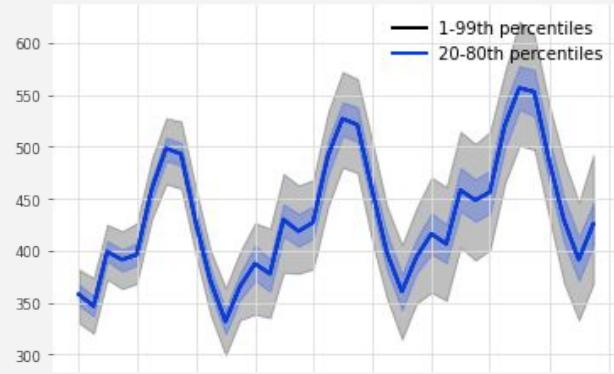
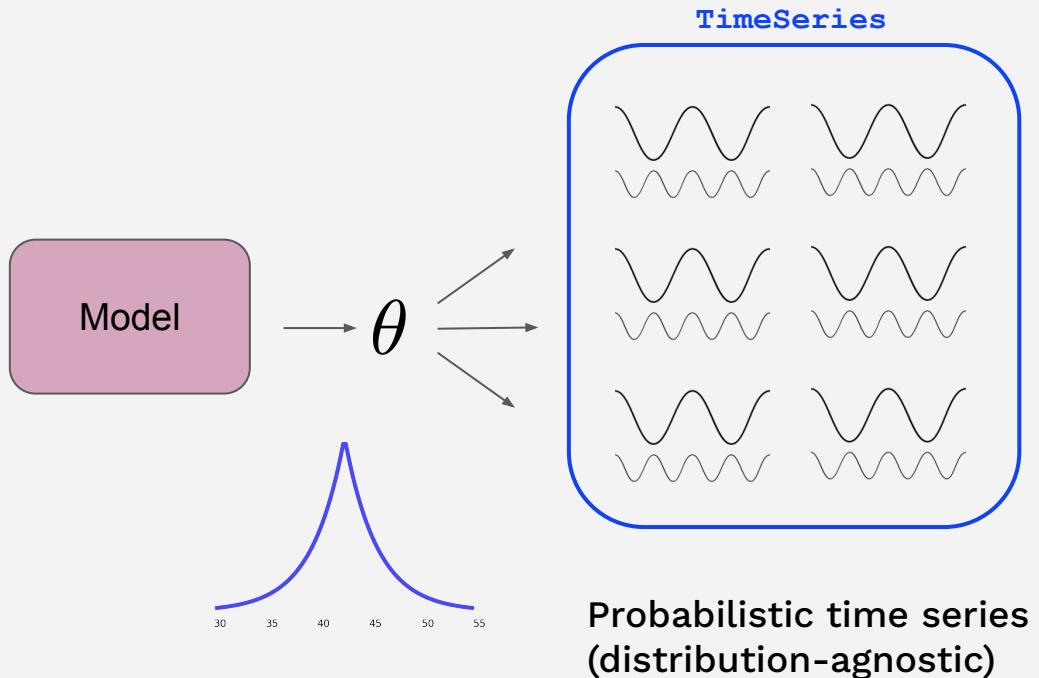
Probabilistic forecasts



A **TimeSeries** contains 3 dimensions



Probabilistic forecasts



● ● ●

```
forecast.plot(low_quantile=0.01, high_quantile=0.99)  
forecast.plot(low_quantile=0.2, high_quantile=0.8)
```

Confidence intervals

Probabilistic forecasts

```
from darts.utils.likelihood_models import (
    BernoulliLikelihood,
    BetaLikelihood,
    CauchyLikelihood,
    ContinuousBernoulliLikelihood,
    DirichletLikelihood,
    ExponentialLikelihood,
    GammaLikelihood,
    GaussianLikelihood,
    GeometricLikelihood,
    GumbelLikelihood,
    HalfNormalLikelihood,
    LaplaceLikelihood,
    LogNormalLikelihood,
    NegativeBinomialLikelihood,
    PoissonLikelihood,
    QuantileRegression,
    WeibullLikelihood,
)
```

- + Priors on distributions' parameters

